

Query-First Semantification of CEUR-WS Proceedings Metadata in Wikidata

Wolfgang Fahl¹[0000–0002–0821–6995], Tim Holzheim^{1,2}[0000–0003–2533–6363],
Christoph Lange^{1,2}[0000–0001–9879–3827], Jerven Bolleman³[0000–0002–7449–1266],
and Stefan Decker^{1,2}[0000–0001–6324–7164]

¹ RWTH Aachen University, Computer Science i5, Aachen, Germany

² Fraunhofer FIT, Sankt Augustin, Germany

³ SIB Swiss Institute of Bioinformatics, Switzerland

Abstract. CEUR Workshop Proceedings (CEUR-WS) has been a non-profit publisher for scientific workshop and conference proceedings in computer science since 1995. For more than 30 years, it has used HTML and PDF as the single source of truth for the metadata required by indexers such as DBLP and the German national library DNB. Scholars have long requested persistent identifiers for the core entities involved: event series, events, proceedings, papers, editors, authors, institutions, and publishers, with the goal of allowing semantic handling and querying with state-of-the-art tools.

In this work, we show how we enabled the semantification using Wikidata as the target knowledge graph infrastructure. The 2015 Semantic Publishing Challenge queries are finally permanently available with up-to-date metadata. The Scholia portal provides user-friendly and aspect-oriented access and is widely used.

We also point out the challenges of aligning the digital, social, and physical aspects of the ongoing transformation and outline a roadmap to use the available queries to bootstrap a new RDF-based metadata-first publishing approach that will mine metadata from proceedings matter and generate the HTML from it instead of the other way around.

Keywords: Knowledge Graph · Linked Data · Metadata Extraction · Publishing · Semantic Web · Wikidata

1 Introduction

Traditional publishing workflows often handle the metadata of the required core entities (cf. Figure 1⁴) in a manner that does not comply with the Findable, Accessible, Interoperable, and Reusable (FAIR) principles⁵, giving low priority to these principles as well as to the Single Source of Truth (SSoT) and Single Point of Truth (SPoT) principles of data management [34]. In this work, SSoT

⁴ blue: main entities in focus for CEUR-WS Wikidata bot

⁵ We refer to such workflows as “unFAIR”; see Section 3.1




Fig. 1. Most relevant entities for scientific publishing (abstract schema).




Fig. 2. SWAT4HCLS2024 example graph walk

refers to the one authoritative source from which metadata originates, whereas SPoT refers to the one authoritative location at which that metadata is accessed by downstream consumers. Therefore, the opportunities to collect the metadata during the lifecycle of an event [20] are missed. As technical editors of CEUR Workshop Proceedings, we have been feeling this pain for over a decade. This work shows how the original HTML/PDF SSoT of CEUR-WS at the ceur-ws.org domain (SPoT) has been used to create a Knowledge Graph representation of the relevant entities in Wikidata between 2022 and 2026 as a step in this direction, thus allowing portals such as Scholia to present the RDF query results in an aspect-oriented and user-friendly way.

1.1 CEUR Workshop Proceedings

Since 1995, CEUR Workshop Proceedings (CEUR-WS)⁶ has published more than 4,200 Volumes and over 90,000 papers as a free online service for publishing proceedings of scientific workshops (and smaller conferences) in computer science, operated by a team of unpaid volunteers.

Technically, all data and metadata of the proceedings are directly represented as a static website in HTML and PDF, served over HTTP and FTP as provided by the workshop editors along the CEUR-WS “how to submit” guide [7].

The metadata is available after some delay of weeks or months through indexing services such as DBLP [27] and K10plus [39]⁷ and at the archive of the TIB <https://www.tib.eu/>.

We report on the successful use of Wikidata and Scholia as the knowledge graph platform to provide FAIR metadata for semantic analysis and querying.

1.2 Contributions of this Paper

The core contributions presented in this work are as follows:

1. To provide the tools, infrastructure, and approaches to consistently supply high-quality FAIR metadata (semantification) for CEUR-WS as detailed in section 3.
2. Splitting the metadata for the three main entities (Event, Event Series, and Proceedings) is a major step towards improving the metadata quality and FAIRness.
3. Providing a bootstrapping approach [3,13] to shift the single point of truth from the current HTML/PDF files to an RDF ready ontology-based format as outlined in Section 6.

2 From Semantic Publishing Challenge 2014 via Text2KG 2023 to Production 2026

The Semantic Publishing Challenge [26,12,36] started in 2014 with Task 1 calling for “answering queries related to the quality of workshops by computing metrics from data extracted from their proceedings, also considering information about persons and events”.

We evaluated the relevance of the original set of 20 queries against three criteria: data availability, coverage, and the need for archival and DOI minting, resulting in a list sorting the queries by priority⁸. The most relevant queries, together with the 5 queries that rely on the main index (Q1.5, Q1.12, Q1.13, Q1.16, and Q1.17), have been taken as the definition of done. We reported on the successful Wikidata export of all CEUR-WS proceeding and event metadata

⁶ <https://ceur-ws.org/>

⁷ <https://dblp.org/>, <https://opac.k10plus.de>

⁸ https://cr.bitplan.com/index.php/List_of_Queries

in 2023 at the TEXT2KG Workshop [20]. The SPARQL queries⁹ are available as named parameterized queries [19].

3 CEUR-WS Semantification in Wikidata

This section describes how the CEUR-WS metadata is extracted from its HTML/PDF single point of truth into the Wikidata knowledge graph. Section 3.1 motivates the semantification from the perspective of the FAIR principles and derives the requirements on queries, ontologies, curation and infrastructure. Section 3.2 then presents the concrete extraction pipeline built around `pyCEURmake`, `Cermine`, `GROBID` and `snapquery`, and the `ceur-spt` prototype that consumes the resulting graph. Figure 3 gives an overview of the overall workflow.




Fig. 3. Workflow of the CEUR-WS Semantification

3.1 Making CEUR-WS FAIR

The FAIR principles [24,42] with Persistent Identifiers (PIDs) and rich metadata as core components are being adopted by an increasing number of publishing organizations.

To make CEUR-WS metadata FAIR, a set of relevant queries should be supported, as derived from the original set of queries of the 2014 Semantic Publishing Challenge as outlined in Section 2. The metadata should be provided based on established ontologies. The manual and automatic curation of entries should be possible with public access for all stakeholders, e.g., editors, authors, organizers, publishers, and indexers. The infrastructure should be stable and there should be sufficient trust in its long-term availability. An open source non-profit infrastructure is preferred since this is also the mode of operation of CEUR-WS.

⁹ https://cr.bitplan.com/index.php/Semantic_Publishing_Challenge_Queries

Given the HTML/PDF/text input of CEUR-WS, the corresponding Knowledge Graph needs to be created and the single-point-of-truth computer-readable metadata has to be separated from the different representations such as HTML so that the above requirements are fulfilled.

Both the HTML and PDF encoding of the original scientific content are structured for the purpose of optimizing the display / output on paper or screens; therefore, there is a structure loss compared to what was originally available in the document source files the authors might have been using [14]. Most scholars are not aware of this loss in the daily use of published content just because the documents are optimized for display and human consumption [44,2]. For the metadata extraction and use in knowledge graphs the difference is sometimes disastrous, e.g., when a simple text from a PDF document cannot be extracted any more due to exotic styling and formatting or simply because only a scanned graphic image version of an older document is available that needs optical character recognition to extract the textual content.

The metadata needed for creating the knowledge graph is only available in natural language/text form and follows rules that have been changed multiple times during the history of CEUR-WS. From 2013 to 2026, there have been 33 different versions¹⁰ of the index file template with no proper tracing of what was changed from version to version. The index and volume HTML files were often edited manually, leading to minor undocumented differences even within these versions. The usage of the different versions follows a long-tail Zipfian distribution with 5 versions covering 60% of all volumes.

3.2 Extracting metadata from HTML and PDF targeting Wikidata

pyCEURmake [17] was created to parse the main HTML index as well as the individual proceedings HTML pages, vastly improving on an earlier simple scraper [29]. The results have been cached in JSON and SQLite storage for quick access. Location entries have been consolidated using the geograpy3 library [35]. Author, Editor and Paper details have been aligned with the DBLP SPARQL endpoint.

Cermine [37] and GROBID [28] have been run on the papers to provide additional metadata references. Figure 4 shows a user-friendly presentation of the data mined by these steps.

For each core entity, we analyzed the availability of Wikidata properties by using our wdgrid analysis tool [16,22], which also generates SPARQL queries. We aligned the queries with the stated needs of the Semantic Publishing Challenge by mapping the wikidata items and properties accordingly.

snapquery [19] was used to make sure that the queries are compatible with different endpoints and do not only run on the official Blazegraph Wikidata query endpoint.

The ceur-spt prototype [18] demonstrates the benefits of semantification by improving linking and navigation capabilities compared to the current static CEUR-WS HTML pages. A complete Cermine run has been integrated so that

¹⁰ https://cr.bitplan.com/index.php/CEUR-WS_Versions

#	Item	Volume	Acronym	Dblp	K10plus	Event	Series
4161	Proceedings of the Workshop on Artificial Intelligence and the Science of Science	Vol-4161	AISSciSci 2025	-	k10plus	Workshop on Artificial Intelligence and the Science of Science	
4160	Proceedings of the International Scientific Workshop on Applied Information Technologies and Artificial Intelligence Systems	Vol-4160	AIT&AIS 2025	-	k10plus	International Scientific Workshop on Applied Information Technologies and Artificial Intelligence Systems	
4159	Proceedings of the 2nd International Workshop on Bioinformatics and Applied Information Technologies for medical purposes (BAITmp 2025)	Vol-4159	BAITmp 2025	-	k10plus	2nd International Workshop on Bioinformatics and Applied Information Technologies for medical purposes (BAITmp 2025)	
4158	Proceedings of the Information Technology and Implementation (ITI&I) Workshop, Artificial Intelligence Technologies and Data Science (IT&I-WS, AITDS 2025)	Vol-4158	ITI&I-WS, AITDS 2025	dblp	k10plus	Information Technology and Implementation (ITI&I) Workshop, Artificial Intelligence Technologies and Data Science (IT&I-WS, AITDS 2025)	
4157	Proceedings of the Workshop on AI for Evidential Reasoning	Vol-4157	AI&EVR 2025	-	k10plus	Workshop on AI for Evidential Reasoning	

(c)2023-2026 Wolfgang Fahl Powered by rickroll

Fig. 4. CEUR Volume Browser wikidata synchronization with downstream dblp and k10plus coverage

paper-level metadata (authors, affiliations, references) is available in the per-volume pages as illustrated for Vol-3262 in Figure 5.

4 Challenges in Aligning the Digital, Social, and Physical Aspects

CEUR-WS is a digital-native, online-only publisher operated by volunteer editors, and serves a scholarly community of editors, authors and research institutions. Downstream indexing activities are handled by DBLP, the German National Library (DNB), TIB, Wikidata/WikiCite and the Scholia community. This section describes the challenges of aligning the digital, social, and physical aspects, since traditional proceedings are still printed on paper and stored on book shelves in libraries.

The three-way alignment is necessary to avoid cognitive dissonance in the discussions and decisions that are needed to support the metadata-first roadmap outlined in Section 8.

4.1 Digital: CEUR-WS as a native online publisher

CEUR-WS has always been online only, without physical artefacts. As discussed in Section 3.1, the current Single Point of Truth encodes metadata in HTML and PDF optimized for printing and human display rather than for machine consumption. There has been a good reason to never run any software such as a triplestore—static content is much more stable and safer than infrastructure that needs maintenance. The Wikidata export makes it possible to query metrics



Fig. 5. Vol-3262 ceur-spt page: structured paper-level metadata derived from the HTML index and PDF mining pipeline (cropped)

such as coverage and consistency. The public synchronization dashboard¹¹ is an example that makes the downstream indexing observable; see Figure 4.

4.2 Social: volunteer ecosystems across institutional boundaries

CEUR-WS editors, DBLP curators, DNB and TIB librarians, Wikidata and WikiCite contributors, and Scholia maintainers all operate as volunteer or mixed-funding communities, each with their own goals and procedures. Coordination across such communities is loosely coupled and based on the trust that what worked in the past will continue to do so in the future. Our work highlights gaps; for example, DBLP is currently indexing only about 80 % of all volumes. Trust in Wikidata as an openly editable knowledge graph has yet to be established, and active checks are needed to detect whether the metadata has been tampered with in unacceptable ways.

4.3 Physical: library catalogues and the shelf model

Library catalogues such as the DNB-curated integrated authority file GND¹² historically optimize for physical books — proceedings that sit on a shelf. For a digital-native publisher there is a mismatch: the distinction between *event*, *event series* and *proceedings*, which the core-entity schema of this paper treats as three separate first-class entities (Figure 1), is not covered adequately in library catalogues. Our effort to convince the DNB to systematically curate event series based on an existing concept in its ontology has so far been unsuccessful.

¹¹ <https://cvb.wikidata.dbis.rwth-aachen.de/wikidatasync>

¹² https://www.dnb.de/EN/Professionell/Standardisierung/GND/gnd_node.html

5 In-Use Impact – WikiCite, Scholia, CEUR-WS

The semantified CEUR-WS metadata is in production use across three overlapping communities: Scholia [31], a portal serving Wikidata content for scholars (Section 5.1); WikiCite [40], the bibliographic database initiative on Wikidata (Section 5.2); and the CEUR-WS stakeholders themselves (Section 5.3). Over 100,000 distinct authors and 6,500 distinct editors with digital traces in the DBLP CEUR-WS index entries span these communities.

5.1 Scholia

Scholia, the most visible in-use consumer of the semantified CEUR-WS metadata, is based on a set of over 380 named parameterized SPARQL queries organized around 40 different aspects. The Wikidata Scholarly Articles Sub-graph Analysis [1] showed that roughly half of the 17 billion triples of the main graph belong to scholarly articles and linked entities. The Wikimedia Foundation needed to perform a graph split due to a 4 TiB maximum storage limit of the Blazegraph backend in use. The Scholia community switched to QLever [4] and made all queries SPARQL 1.1 compliant [43]. Note that, before the graph split, there was a hold on adding more scholarly articles to Wikidata, which is why CEUR-WS paper metadata has not yet been mass-inserted. Writes to Wikidata are performed through the officially approved bot *User:CEUR-WS* [8], for which the community permission was obtained in 2023 [41].

Table 1 shows 12 SPARQL-driven panels on the Scholia event aspect [31]. All queries are enabled by this work and rendered identically by the legacy Blazegraph endpoint at <https://scholia.toolforge.org/event/Q124669329> and by the post graph-split QLever endpoint [4,43] at <https://qllever.scholia.wiki/event/Q124669329>, illustrated for SWAT4HCLS 2024 (Q124669329)¹³. All 12 panels survived the 2025 Wikidata graph split unchanged; only the underlying engine and the SPARQL version (1.1 compliant) differ.

Figure 6 shows the user-friendly visualization of the coauthorship graph.

5.2 WikiCite

WikiCite [40] provides the bibliographic backbone on which CEUR-WS metadata is contributed to Wikidata.

Figure 7 shows the distribution of Wikidata editors who have created or revised CEUR-WS *proceedings* items (conference proceedings (Q1143604)¹⁴). The long-tail distribution illustrates that while our bot accounts for the bulk of systematic mass contributions, a broad community of individual WikiCite editors continuously curates, corrects and enriches the proceedings metadata. This overlap of bot-driven bulk ingestion and human community curation is a core strength of the WikiCite approach and a prerequisite for the Scholia panels

¹³ <https://www.wikidata.org/wiki/Q124669329>

¹⁴ <https://www.wikidata.org/wiki/Q1143604>

Table 1. Scholia SWAT4HCLS 2024 panels based on SPARQL queries

#	Panel/Query	style	entities
1	data	header	event
2	people	table	scholar
3	co-authors	graph	scholar
4	sponsors	table	institution
5	timeline	timeline	event
6	topic-scores	table	
7	uses	table	
8	recent-publications	table	paper, scholar
9	proceedings	table	proceedings, paper
10	presentations	table	
11	related-events-timelocation	table	event
12	related-events-people	table	event, scholar

discussed above. The plot was generated with `ceur-wd-contributions` [17] by traversing the revision history of all CEUR-WS proceedings items on Wikidata.

5.3 CEUR-WS stakeholders

The bot currently exports *proceedings* and *event* items only. Items for the related entities event series, paper, scholar and institution are created and edited by the CEUR-WS stakeholders active in the Wikidata community. The SWAT4HCLS 2024 case has already been presented in Section 1 (graph walk, Figure 2) and in this Section 5 (Scholia panels, Table 1 and Figure 6).

6 Single Point of Truth Bootstrapping

The work reported so far uses the CEUR-WS HTML/PDF artefacts as the single point of truth (SPoT) and derives the Wikidata knowledge graph from them. This section outlines how that relationship is being inverted in a bootstrapping fashion [3,13], so that a machine-readable, ontology-based representation becomes the SPoT and the HTML/PDF artefacts are generated from it. Section 6.1 introduces the CEUR-WS Proceedings Ontology.

6.1 CEUR-WS Proceedings Ontology

The CEUR-WS Proceedings Ontology¹⁵ formalizes the structure of the semantified data. It covers all entity classes defined by the Semantic Publishing Challenge (see Figure 8). To maximize reusability and interoperability, terms from

¹⁵ <https://github.com/ceurws/semantify/blob/main/ontology/ceur-ws.ttl>




Fig. 7. Editor activity distribution across 4,174 CEUR-WS proceedings items on Wikidata, showing the complementary roles of automated bulk ingestion and targeted human curation of the WikiCite community (55 distinct editors, 21,405 total contributions).

7 Lessons Learned

The effort required for this work was substantially higher than initially anticipated, as a large number of non-trivial edge cases in the legacy metadata had to be resolved individually and manually. The following lessons generalize beyond CEUR-WS.

7.1 Queries as first-class citizens

The Semantic Publishing Challenge and the Scholia portal both use queries as the core structuring elements. Using the guiding principle “If you cannot query it, it does not exist” changes the view on how to provide metadata. While traditionally schemas and ontologies are created first to allow querying for use-case specific areas of interest, the query-first and query as first-class citizens approach describes the information needs of stakeholders as queries, similar to competency questions [38] in ontology engineering. Having the queries as named, parameterized, technology-independent entities themselves is helpful [19]. Nielsen’s Synia [30] approach shows a lightweight application of this idea.

A main success factor for this work was being able to apply metadata to the queries. This way the queries could be ordered by task and/or relevance and linked to their respective implementation. These links were just motivations – as first-class citizens the queries are available with proper URIs as persistent identifiers and part of a meta-knowledge graph that is useful for the further semantification steps.

2. CEUR-WS Proceedings Ontology: Overview

[back to ToC](#)

This ontology has the following classes and properties.

Classes

Academic Event	Agent	Authorship	Bibliographic Resource	Conference	Contribution	Editorship	Event Series	Grant	Organization	Paper	Person
Proceedings Volume	Session	Workshop									

Object Properties

affiliated with	associated event	co-located with	contains paper	contributor	contributor of	coordinates	funded by	funding agency	has contributor
has session	part of series	presented at	programme committee chair						

Data Properties

acronym	article type	author name string	call for papers URL	citation type	country	DBLP author ID	DBLP event ID	DBLP publication ID	DBLP venue ID
DOI	EU project ID	full work URL	GND ID	grant ID	introduces ontology	K10plus ID	mentions ontology	name as listed	official website
ROR ID	session order	short name	URN-NBN	volume number	WikiCFP ID	WikiCFP series ID	Wikidata ID	ORCID	ordinal

Named Individuals

0000_0003_2533_6363	Book	Book Chapter	CEUR Workshop Proceedings	CEUR-WS Article Type Scheme	CEUR-WS Citation Type Scheme	CEUR-WS.org
Conference Paper	Demo Description	Extended Abstract	Full Paper	Invited Paper	Journal Article	Position Paper
System Description	Technical Report					

Fig. 8. Overview of the proposed CEUR-WS Proceedings ontology

7.2 Property creation as a community process

Wikidata properties such as *colocated with* (P11633)¹⁷ undergo a Wikidata’s property proposal process¹⁸ before a new property is created. As described in [20] the creation of this property proposal¹⁹ took some effort and led to the clarification that the property is asymmetric. The property was considered as highly relevant and well defined.

7.3 Legal aspects of Publishing Personal Data of Scholars

The proposed workflow calls for publishing personal data of authors and editors early in the life cycle of an event. According to the legal advisors we consulted, this is legal since a combination of conditions (Art. 6 GDPR) is met: Explicit consent is given by the stakeholders (Art. 6(1)(a)), the legitimate interest of the stakeholders is pursued (Art. 6(1)(f)), the data might be public already, the processing is necessary for scientific research (Art. 6(1)(e), Art. 89). The strongest condition is that there might be a legal obligation (Art. 6(1)(c)) as would be the case when publishing via major print publishers, which must provide a copy of the proceedings to the German National Library under German law. Consent, Public Data, and Contractual Necessity are crucial for scholarly publishing. The legal basis and data-subject rights for publication of personal data of scholars by CEUR-WS are documented in the CEUR-WS privacy policy²⁰ and tracked in the project issue tracker issue #11 of the CEUR semantify project²¹.

¹⁷ <https://www.wikidata.org/wiki/Property:P11633>¹⁸ https://www.wikidata.org/wiki/Wikidata:Property_proposal¹⁹ https://www.wikidata.org/wiki/Wikidata:Property_proposal/colocated_with²⁰ <https://ceur-ws.org/privacy-dsgvo.html>²¹ <https://github.com/ceurws/semantify/issues/11>




Fig. 9. Year-split prototype of the CEUR-WS index (`index_alt.html`): a step towards iterative replacement of the monolithic index page.

7.4 Incremental migration over one-shot replacement

The challenge in introducing the new CEUR-WS workflow is to find a sweet spot between effort and benefit [33]. Given the long-tail distribution of ever more exotic corner cases that have shown up in the analysis of the existing data and workflow there is potentially a huge list of “small” problems that would require an inadequately high effort to be solved systematically. The “sweet spot” is in solving only the most relevant problems and not fixing some of the exotic details at all or doing it manually. Therefore we will avoid a “big bang” introduction, which would replace the old approach in one single big (and therefore very risky) step. Instead, we propose to make smaller changes gradually. As a first step, the main *index.html* shall be split by years (with a link to the old legacy complete list). Then, the per-year list may be generated from the single point of truth data. This will be continued iteratively, going from the current year back to a point where it seems no longer reasonable to spend further effort.

A first bootstrapping prototype of this year-split is available (see Figure 9).

8 Roadmap: Metadata-First Publishing

The coverage of the core entities Event Series, Papers, Scholars and Institutions needs to be improved by consolidating and completing the RDF data of the downstream indexing services at DBLP and German National Library (DNB). The central long-term objective of this effort is to shift the CEUR-WS single point of truth from the current HTML/PDF artefacts to a workflow in which the metadata is mined from the proceedings matter and the rendered HTML and computer readable RESTful content-negotiation artefacts should be generated from it as static content. The 2026 migration of the CEUR-WS infrastructure from RWTH Aachen University to Gesellschaft für Informatik provides

the opportunity to take further concrete steps in this direction. The open-source CEUR-WS semantify project [23] coordinates this work.

8.1 Open community proposals to be integrated into the future workflow

The move from the historical `ceur-make` / `ceur-make-gui` toolchain to a future `proceedings-matter` → `SPT` → `HTML` workflow is not yet decided. Several complementary proposals have been contributed during the timeframe of this work; each addresses part of the future workflow, and a central open task of the roadmap is to integrate them into a coherent pipeline:

- *Author-driven metadata entry*: Michael Cochez proposes re-using the `ceur-ws-prototype` [11], inspired by `ceur-make-ui` [6], with a simple HTML form so that authors enter their own metadata, thus automating volume creation and reducing editor workload.
- *Automated extraction and compliance checking*: Jerven Bolleman proposes automating metadata/PDF extraction, index generation and CEUR compliance checks using `ceur-index-html-helper` [5], already applied to generate <https://www.jerven.eu/swat4hcls-proceedings-2026/>.
- *SPT-first research prototype*: the authors propose making the `ceur-spt` backend fully functional first and only then adding the frontend, following the `SemPubFlow` prototype approach.
- *Historical context*: the 2019 XML-based proposal by Sarven Capadisli and the ESWC 2022 background discussion by Nathanael Arndt are acknowledged as previous ideas and work.

9 Conclusion

CEUR-WS proceedings metadata is systematically made available in Wikidata via the CEUR-WS bot. The relevant Task 1 Queries of the Semantic Publishing Challenge of 2014/2015 are now available on different endpoints as named parameterized, i.e., FAIR queries, namely the SPARQL 1.1-compatible QLever endpoint. The Scholia portal provides production access to queries based on and extending this set. The Wikidata, WikiCite, Scholia, and CEUR-WS stakeholder communities have appreciated and contributed to this work.

The contributions of this work are valuable for other publishing use cases²² that might struggle with the same pain points.

9.1 Limitations

The contributions reported here are subject to several limitations. First, Wikidata coverage is currently complete for the *proceedings* and *event* entities only; items for papers, scholars, institutions and event series rely on complementary

²² <https://cr.bitplan.com/index.php/PublishingVenues>

community curation and are therefore not uniformly complete. Second, the reconciliation pipeline depends on the availability and quality of external endpoints (DBLP, K10plus, geograpy3); degradation or schema changes in these sources propagate into the pipeline. Third, roughly 20 % of CEUR-WS volumes are not yet indexed by DBLP, which limits editor and author disambiguation for that subset. Fourth, the current Single Point of Truth remains the legacy HTML/PDF representation; the metadata-first workflow outlined in Section 8 has not yet replaced it in production. Finally, generalisation beyond CEUR-WS to other publishers has not been empirically evaluated and is left for future work.

9.2 Supplemental Material Statement

Source code for the prototypes and semantification tools is available via the repositories of the github ceurws organization²³. Extended supplemental material — including the schema diagrams (Figures 1 and 2) in their editable Graphviz form, the SPARQL queries derived from the Semantic Publishing Challenge, and background research — is provided by the CR research wiki [15]. The LaTeX sources, bibliography, build pipeline and the `AGENTS.md` governance policy for AI-assisted editing of this paper are available at [21]. The cited software artefacts are released under Apache-2.0: `ceur-spt` [18], `pyCEURmake` [17], `snapquery` [19], and `wdgrid` [16].

9.3 Acknowledgements

We would like to thank Jakob Voß for helping with the k10plus matching and creating the Wikidata property colocated with (P11633)²⁴ in due time.

Thomas Hoeren and Jonas Kuitert of ITM, Münster kindly gave hints for legal aspects as summarized in section 7.3.

This paper is dedicated to the memory of CEUR-WS board member Ralf Klammer † January 2023.

This research has been partly funded by a grant of the German Research Foundation (DFG).²⁵

9.4 Declaration of use of Generative AI

The authors used AI-based coding assistants during the engineering and preparation of this work. Use of such assistants is governed by an `AGENTS.md` policy file maintained in this paper’s repository and, equivalently, in the source repositories of the software artefacts `ceur-spt` [18] and `pyCEURmake` [17]. The policy mandates:

²³ <https://github.com/ceurws>

²⁴ <https://www.wikidata.org/wiki/Property:P11633>

²⁵ ConFIdeNT project; see <https://gepris.dfg.de/gepris/projekt/426477583>, National Research Data Infrastructure (NFDI) for Data Science and Artificial Intelligence; see <https://gepris.dfg.de/gepris/projekt/460234259>

- a plan-and-confirm protocol: the assistant must state its analysis, plan, and scope, and receive explicit human confirmation before any read, write, or execution;
- a strict scope for the manuscript itself: AI assistance is limited to spelling/grammar corrections and formatting support (tables, figures, build automation). AI was not used to generate, rewrite or contribute original scientific content, arguments, or prose;

All scientific claims, interpretations and conclusions are the sole responsibility of the authors.

References

1. AKhatun: Wikidata Scholarly Articles Subgraph Analysis, https://wikitech.wikimedia.org/wiki/User:AKhatun/Wikidata_Scholarly_Articles_Subgraph_Analysis
2. Ansoerge, L.: Let’s publish full-text scientific articles in html, not just pdf. *European Science Editing* **47** (Nov 2021). <https://doi.org/10.3897/ese.2021.e75834>
3. Bardini, T.: Bootstrapping: Douglas Engelbart, Coevolution, and the Origins of Personal Computing. Stanford University Press, USA (2001)
4. Bast, H., Buchhold, B.: QLever. In: *Proceedings of the 2017 ACM Conference on Information and Knowledge Management*. pp. 647–656. ACM (Nov 2017). <https://doi.org/10.1145/3132847.3132921>
5. Bolleman, J.: ceur-index-html-helper. <https://github.com/JervenBolleman/ceur-index-html-helper>
6. CEUR-WS community: ceur-make-ui. <https://github.com/ceurws/ceur-make-ui>
7. CEUR-WS Editor Team: How to submit to ceur workshop proceedings. <https://ceur-ws.org/HOWTOSUBMIT.html> (2019)
8. CEUR-WS Editor Team: Wikidata bot account: User:CEUR-WS (2023), <https://www.wikidata.org/wiki/User:CEUR-WS>
9. CEUR-WS semantify community: Sample volume subgraphs (discussion #15). <https://github.com/ceurws/semantify/discussions/15> (2026)
10. CEUR-WS semantify community: Schema consolidation (discussion #6). <https://github.com/ceurws/semantify/discussions/6> (2026)
11. Cochez, M., contributors: ceur-ws-prototype. <https://github.com/ceurws/ceur-ws-prototype>
12. Di Iorio, A., Lange, C., Dimou, A., Vahdati, S.: Semantic publishing challenge – assessing the quality of scientific output by information extraction and interlinking. In: Gandon, F., Cabrio, E., Stankovic, M., Zimmermann, A. (eds.) *Semantic Web Evaluation Challenges*. pp. 65–80. Springer International Publishing, Cham (2015)
13. Engelbart, C.: About Bootstrapping. <https://www.dougenelbart.org/content/view/226/269/> (2007)
14. Fahl, W.: The history of scientific publishing (2023), https://cr.bitplan.com/index.php/The_History_of_Scientific_Publishing
15. Fahl, W.: CR: Semantic MediaWiki on semantification, knowledge graphs and scientific publishing. https://cr.bitplan.com/index.php/Main_Page (2026)
16. Fahl, W.: wdgrid: NiceGUI-based Wikidata grid display and sync (2026). <https://doi.org/10.5281/zenodo.20035246>

17. Fahl, W., Holzheim, T., Ayan9801: pyCEURmake: Python implementation for CEUR Workshop Proceedings (2026). <https://doi.org/10.5281/zenodo.20034998>, also known as cvb (CEUR Volume Browser)
18. Fahl, W., Holzheim, T., byzakyz: ceur-spt: RESTful Single Point of Truth server for CEUR-WS (2026). <https://doi.org/10.5281/zenodo.20037981>
19. Fahl, W., Holzheim, T., Hoyt, C., Priskorn, D., et al.: SnapQuery: FAIR Management of Query Sets to Mitigate Query Rot in Knowledge Graphs (2026). <https://doi.org/10.5281/zenodo.17802424>
20. Fahl, W., Holzheim, T., Lange, C., Decker, S.: Semantification of CEUR-WS with wikidata as a target knowledge graph. In: Joint Proceedings of Text2KG and BiKE co-located with ESWC 2023. CEUR Workshop Proceedings, vol. 3447, pp. 205–225. CEUR-WS.org (2023), https://ceur-ws.org/Vol-3447/Text2KG_Paper_13.pdf
21. Fahl, W., Holzheim, T., Lange, C., Decker, S.: CEURWS_Semantification_ISWC2026: Paper source repository with AGENTS.md governance policy. https://github.com/WolfgangFahl/CEURWS_Semantification_ISWC2026-paper (2026)
22. Fahl, W., Holzheim, T., Westerinen, A., Lange, C., et al.: Property cardinality analysis to extract truly tabular query results from wikidata. In: Proceedings of the 3rd Wikidata Workshop 2022. CEUR Workshop Proceedings, vol. 3262. CEUR-WS.org (2022), <https://ceur-ws.org/Vol-3262/paper7.pdf>
23. Fahl, W., et al.: semantify: Semantification of the ceur-ws publication workflow. <https://github.com/ceurws/semantify> (2026)
24. GO FAIR International Support and Coordination Office: Fair principles (2019), <https://www.go-fair.org/fair-principles/>
25. Knublauch, H., Kontokostas, D.: Shapes constraint language (SHACL). Tech. rep., W3C (Jul 2017), <https://www.w3.org/TR/shacl/>
26. Lange, C., Di Iorio, A.: Semantic publishing challenge – assessing the quality of scientific output. In: Communications in Computer and Information Science, pp. 61–76. Springer International Publishing (2014). https://doi.org/10.1007/978-3-319-12024-9_8
27. Ley, M.: DBLP. Proceedings of the VLDB Endowment **2**(2), 1493–1500 (Aug 2009). <https://doi.org/10.14778/1687553.1687577>
28. Lopez, P.: GROBID: Combining automatic bibliographic data recognition and term extraction for scholarship publications. In: Research and Advanced Technology for Digital Libraries, pp. 473–474. Springer Berlin Heidelberg (2009). https://doi.org/10.1007/978-3-642-04346-8_62
29. Nielsen, F.Å.: Scraping ceur workshop proceedings, <https://github.com/WDscholia/scholia/blob/master/scholia/scrape/ceurws.py>
30. Nielsen, F.Å.: Synia: Displaying data from wikibases (2023). <https://doi.org/10.48550/ARXIV.2303.15133>
31. Nielsen, F.Å., Mietchen, D., Willighagen, E.: Scholia, scientometrics and wikidata. In: Blomqvist, E., Hose, K., Paulheim, H., Ławrynowicz, A., et al. (eds.) The Semantic Web: ESWC 2017 Satellite Events. pp. 237–259. Springer International Publishing, Cham (2017). https://doi.org/10.1007/978-3-319-70407-4_36
32. Peroni, S., Shotton, D.: The spar ontologies. In: The Semantic Web – ISWC 2018. pp. 119–136. Springer International Publishing, Cham (2018)
33. Pohl, K., Rupp, C.: Requirements engineering fundamentals. Rocky Nook, Santa Barbara, CA, 2 edn. (May 2015)
34. Queiroz, M., Tallon, P., Coltman, T.: Data value and the search for a single source of truth: What is it and why does it matter? In: Proceedings of the 57th

- Hawaii International Conference on System Sciences. HICSS, Hawaii International Conference on System Sciences (2024). <https://doi.org/10.24251/hicss.2024.795>, <http://dx.doi.org/10.24251/HICSS.2024.795>
35. Rakshit, S., et al.: geograpy3: Extract place names from text and resolve them to countries/regions. <https://github.com/somnathrakshit/geograpy3> (2024)
 36. Ronzano, F., del Bosque, G.C., Saggion, H.: Semantify CEUR-WS Proceedings: Towards the Automatic Generation of Highly Descriptive Scholarly Publishing Linked Datasets, pp. 83–88. Springer International Publishing (2014). https://doi.org/10.1007/978-3-319-12024-9_10
 37. Tkaczyk, D., Szostek, P., Fedoryszak, M., Dendek, P.J., et al.: CERMINE: automatic extraction of structured metadata from scientific literature. International Journal on Document Analysis and Recognition (IJ DAR) **18**(4), 317–335 (Jul 2015). <https://doi.org/10.1007/s10032-015-0249-8>
 38. Uschold, M., Gruninger, M.: Ontologies: principles, methods and applications. Knowl. Eng. Rev. **11**(2), 93–136 (Jun 1996)
 39. Wiermann, B.: K10plus – Zehn Bundesländer in einem Bibliothekssystem (2019), <https://blog.slub-dresden.de/beitrag/2019/03/27/k10plus-zehn-bundeslaender-in-einem-bibliothekssystem>
 40. WikiCite Contributors: WikiCite: A bibliographic database based on Wikidata. <http://wikicite.org/>
 41. Wikidata Contributors: Requests for permissions/bot/ceur-ws (2023), https://www.wikidata.org/wiki/Wikidata:Requests_for_permissions/Bot/CEUR-WS
 42. Wilkinson, M.D., Dumontier, M., Aalbersberg, I.J., Appleton, G., et al.: The FAIR guiding principles for scientific data management and stewardship. Scientific Data **3**(1) (Mar 2016). <https://doi.org/10.1038/sdata.2016.18>
 43. Willighagen, E.L., Mietchen, D., Patel-Schneider, P., Linden, K., et al.: Scholia 2026: compliance with SPARQL 1.1. <https://raw.githubusercontent.com/WolfgangFahl/ScholiaGraphSplitPaper/main/main.pdf> (2026), accepted for SWAT4HCLS 2026; CEUR-WS volume forthcoming (ISSN 1613-0073)
 44. Yu, C., Zhang, C., Wang, J.: Extracting body text from academic PDF documents for text mining. CoRR **abs/2010.12647** (2020), <https://arxiv.org/abs/2010.12647>